

# A Doomsday Vault of Software Engineering Tools

Archiving Software Engineering Tools from ICSE and FSE 2011 through 2014

Emerson Murphy-Hill  
Department of Computer  
Science  
NC State University  
Raleigh, NC, USA 27695  
emerson@csc.ncsu.edu

## ABSTRACT

Many innovative software engineering tools appear at the field's premier venues, the International Software Engineering Conference (ICSE) and the Foundations of Software Engineering (FSE). But what happens to these tools after they were presented? In this paper, we spend 10,000 hours trying to obtain, download, use, and repackage 150 tools from ICSE and FSE's tool demonstration tracks. Our results enumerate the practical and accidental reasons that software engineering tools fail to work over time, and provide practical implications for creating lasting tools.

## CCS Concepts

•General and reference → *Empirical studies*;

## Keywords

Software engineering tools; replication

## 1. INTRODUCTION

Software engineering research seeks to better understand software and how it is built and constructed. While such understanding in isolation can be insightful, ultimately a substantial amount of such research aims to impact the practice of software engineering. To do so, researchers can create recommendations for new software engineering practices, can create educational techniques and materials, and can create tools.

Arguably creating new tools is the most common way that software engineering researchers attempt to influence practice. Broadly speaking, tools are software that can help design and build software. Examples include a tool that helps mobile application developers choose which devices to target [9], a tool that checks the use of locks in multithreaded programs [4], and a tool that creates code from natural language text [3].

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than ACM must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from [permissions@acm.org](mailto:permissions@acm.org).

FSE '16 November 13–19, 2016, Seattle, WA, USA

© 2016 ACM. ISBN 978-1-4503-2138-9.

DOI: 10.1145/1235

While papers describe tools in software engineering venues, there are several reasons why software engineering researchers should make the tools themselves available. First, the full details of how the tool works, both in terms of its internals and from an end-user perspective, may not be clear from the paper. Second, making the tools available can help practitioners try and adopt the tools in their work. Third, it helps facilitate reproducibility by enabling future researchers to perform studies with the original tools. Finally, it helps the advancement of the field by allowing others to build on existing tools, rather than re-building them from scratch.

Despite the benefits of making tools available, in this paper we catalog the practical difficulties in doing so. We describe a study that examined the tools presented at the premier venues for software engineering research over the past several years. As part of a graduate university course on software engineering, we spent 10,000 hours trying to obtain, download, use and repackage tools from the International Conference on Software Engineering (ICSE) and the Symposium on the Foundations of Software Engineering (FSE). In doing so, we make the following three main contributions in this paper:

- A study that evaluates the difficulty in getting tools working across a variety of software engineering research.
- A synergistic course project for graduate software engineering students that gives students meaningful educational value *and* provides the research community value.
- XXX existing tools repackaged in automatically-built virtual machines to make it easier for others to use these tools.

## 2. RELATED WORK

timeliness

General scientific interest in reproducibility. How it fits with other Rs. State of practice.

Software reproducibility outside SE. Systems [2, 1]

Others [8, 13, 12, 7].

Arguably as SE researchers we should be the best!

Other fields: - Data science (doing better - <http://zenodo.org>)

- MPC journal requires code submission (tarball) - In general, Math not doing a lot with code sub submissions - Security, should be able to do VM, but don't (but see Will's

paper) - HPC doesn't (and maybe can't) - RTC is starting to do it, but maybe shouldn't

"We argue that, with some exceptions, anything less than the release of source programs is intolerable for results that depend on computation." [6].

Constant worry is protecting IP in security and HPC

Repeatability in AI-SE: "depends on who does it" Reproducibility assessment for 2 papers, plus reproducibility overview [5]

While no studies of software reproducibility in our field, practical interest.

Artifact evaluation committees. (SIGPLAN, where else?)

### 3. RESEARCH QUESTIONS

Henceforth, we will simply say *tool* to refer to software engineering tools presented at the International Conference on Software Engineering or Foundations of Software Engineering in their respective tool demonstration tracks.

1. How much effort is required to get tools to work?
2. What are the barriers to get tools to work?
3. How much effort is required to get tools to work in virtual machines?
4. What are the barriers to get tools to work in virtual machines?
5. To what extent is this study practical to implement in a classroom context?

### 4. COURSE DESCRIPTION

We conducted the study described in this paper as part of a graduate software engineering course in the Computer Science department at North Carolina State University. The department has about 200 PhD students and 450 MS students in its graduate program. While graduate students are not required to take the course, it is one of seven core systems courses, from which students must take one course. The department does have a specialty MS degree track in software engineering,<sup>1</sup> which does require this course.

Course content covers software engineering processes, software architecture, design patterns, software security, verification and validation, estimation, project management, requirements, certification, and formal methods.<sup>2</sup> Apart from the course content, involving lectures and a final exam, the other major component of the course is the project. Next, we describe the project for this course in prior offerings of the course (Section ??) and in the offering described in this paper (Section ??).

#### 4.1 Prior Course Project and Criticism

In prior offerings of the course, students were essentially given two project offerings. In one version, students could create a software engineering tool, such as a plugin for Eclipse. In the other version, students chose several existing, similar software engineering tools, and applied them to open source software, then reflected on what they learned about the tools

<sup>1</sup><https://www.csc.ncsu.edu/academics/graduate/degrees/se.php>

<sup>2</sup>Syllabus: TODO

and the projects. Most students opted for the second version. In both cases, students were required to write up their results. This project is similar to a course project assigned by David Notkin at the University of Washington.<sup>3</sup>

After offering this project to students for several years, the instructor recognized several problems the way he had executed it:

- The course did not have enough time to teach students technical writing skills, and thus final papers were of poor quality, on average. Moreover, while technical communication is a course objective, other types of communication would likely be more valuable to most students, who are non-thesis and industry-focused.
- Students had limited exposure to state-of-the-art tools; the tools they chose to study were typically quite basic.
- Students appeared to rarely gain technical skills during the project.
- Most students were not doing *novel* work; each semester, different students wrote similar reports on similar tools, which had no value beyond their educational value.

While the last point may seem odd – why would a course project need value beyond its educational value? – it's not unusual for some software engineering courses to have added value, such as contributing to open source projects [11, 10]. In the case of the present course, the instructor thought that such system development would not be useful to most students in the course, who typically come in with 1 to 2 years of industrial software engineering experience.

#### 4.2 New Course Project

In Fall of 2015, the instructor changed the course project to alleviate the problems described in the last subsection. In short, student teams were assigned tools described in a prior research paper at ICSE or FSE, the premier venues for software engineering research. Students were required to obtain the tools, get them running, and redistribute them.

The instructor chose tool demonstration papers, rather than full technical papers, for practical reasons. Full technical papers may not present a tool; for instance, a purely qualitative study that reports on empirical findings may have no software to go along with it. In contrast, tool demo papers almost certainly had a working tool when the paper was presented at a conference.

There were several learning goals of the course project:

- Gain deep experience with several state-of-the-art software engineering tools, and broad overview of many others;
- Effectively read research papers;
- How to build virtual machines that contain custom software;
- How to script virtual machine creation; and
- Oral communication skills.

In the remainder of this section, we describe project activities, requirements, and deliverables.

<sup>3</sup><https://courses.cs.washington.edu/courses/cse503/>

#### 4.2.1 Team Formation and Tool Selection

The instructor formed teams of about 5 students by randomly selecting about 4 on-campus students and possibly 1 off-campus, distance education student.

During class, the instructor asked teams to find the ICSE and FSE demonstration tracks online, skim the papers from those tracks, and extract several pieces of information: the venue, the paper name, the tool name, any links to source code or binaries for the tools, the technologies involved in the tool's creation, and a rough estimate of how difficult the team thought it would be to get the tool working. Teams put this information on a shared spreadsheet with one row per paper. We collected papers from from 2014, the most recent year either ICSE or FSE papers will officially posted online, to 2009. In total, we collected 188 papers from 12 conferences.

Teams were then given the opportunity to look over the list of tools, and identify tools they might want to work on. During the following class period, teams chose  $N$  tools to work on, where  $N$  was the number of people on the team. Each tool could be assigned to one and only one tool.

Because some tools were in high demand and some in low demand, for fairness, teams chose tools in a round-robin draft, that is, one team chose a tool, then the next team chose a tool, and so on, until all teams had chosen one tool, and then the first team chooses a second tool, and the next team chooses a second tool, and so on. At the time of the draft, 100 students in the class, so the tools eligible for selection were the 100 most recent tools, which included all tools from 2012–2014, and a few from 2011.

#### 4.2.2 Obtaining Tools

After teams were assigned tools, their first task was to obtain the tools, including the source code, if possible, and the binary if not. For tools that were deployed (or partially deployed) via a web service, students were required to obtain the source or binary for the service as well; in short, the students needed to obtain all software components necessary to run the tool in isolation.

Teams were instructed to first try to obtain the tool from links in the paper itself or via the web. If teams could not find the tool, the teams were instructed to ask the authors, using an email template show in Figure 1. To minimize communication with the authors, the email contained requests for several author piece of information, the need for which will become clear later in this section. Students were instructed to customize this template, but to use it as a starting point. Full instructions for use of this template can be found online.<sup>4</sup>

#### 4.2.3 Getting the Tools Working

Once students had obtained the tools, they had two weeks to get the tools working. Because some tools may only work in the very specific situations described in their papers, teams were required to get the tools “working” as it was described in the paper. Teams were allowed to use any means necessary to get the tools working, including asking classmates for help, but were discouraged from harassing the tools’ authors. At the end of the two weeks, if the tool

To: <AUTHORS OF PAPER, BUT NOT MORE THAN THE FIRST 5 AUTHORS>

Subject: Regarding your tool, <TOOLNAME>

Dear Dr. <LASTNAME> and colleagues,

I enjoyed reading your paper <PAPER TITLE>. As part of a graduate software engineering class at NC State University, my team has chosen to use the tool <TOOLNAME> as part of our class project. In short, the class is using tools from the past few ICSEs and FSEs, then putting all those tools in an accessible form (e.g., virtual machine) in central location. Our project is supervised by Dr. Emerson Murphy-Hill (CC'd via ncsu-csc-510googlegroups.com). When our project is complete, we plan on aggregating our results (e.g., how many tools could we get working, how easy was it, etc) into a research paper.

Because my grade rests on your tool, I am motivated to get it working. I have a few questions for you before I get started.

May I have permission to redistribute an executable version of your tool? Specifically, we plan on putting your tool in a virtual machine, then posting that VM on the public internet.

May I have permission to redistribute the source code for your tool? Specifically, I plan on getting the tool building into the virtual machine (by way of Vagrant) and posting the build scripts on GitHub.

Could you please send me a link (or attachment) to the executable version of your tool? I have attempted to find the tool on the internet, but have been unsuccessful.

Could you please send me a link (or attachment) to the source code of your tool? I have attempted to find it on the internet, but have been unsuccessful.

I've had some trouble getting the tool to work... <describe what you tried, describe what you expected to happen, describe what actually happened.>

Thank you very much,

<YOUR NAME>

**Figure 1: An email template used by students for obtaining tools.**

was obtained and working, teams certified that the tool as working on the spreadsheet.

If a team could not get a tool working for whatever reason, several things happened. First, the team must certify the tool as “unworkable.” Second, the team would be assigned a new tool by moving down the list of tools, and the process would start again. Third, the tool certified as “unworkable” would be made available to the other teams to try to get working.

The instructor created a disincentive for students to unnecessarily certify a tool as unworkable. If a team got an unworkable tool working, the team that certified it as unworkable had their final course grade reduced by a minor grade (for example, from an A to an A- or from a C to a C-). The team that gets the “unworkable” tool working gets their final grade increased by a minor grade point. Furthermore, the only way to get an A+ course is to be on a team that gets an unworkable tool working.

#### 4.2.4 Getting the Tools Working in a Virtual Machine

The first graded deliverable was a VirtualBox virtual machine image that contained an operating system, the tool,

<sup>4</sup><https://docs.google.com/document/d/1dWvHS8gf37MgYafqOSl3BJzHYkdG1BHgDT4CYe-QN1I/edit?usp=sharing>

any and other software required by the tool (such as Eclipse), documentation, and license information. The goal of creating this image was to make it as easy as possible for future potential users to try the tool out; in essence, any “fiddling” required to get a tool working would not have to be done by the user.

The grading rubric for the image<sup>5</sup> additionally specified that minimal work is required on the part of the user to see the tool in action, that the technology stack does not contain any proprietary software that is not strictly necessary, and that the image is as small as possible.

#### 4.2.5 *Building the Virtual Machine Image Automatically*

After building a basic virtual machine image by hand, teams next built Vagrant scripts that built virtual machines automatically. The main goal of creating the script was to add transparency to the process of installing the tool; if future users want to install the tool in their own development environment, the script provides a specification for doing so.

The grading rubric for the virtual machine script<sup>6</sup> was largely the same as for the hand-built image. For instance, some things that were easy to do in a hand-built image turned out difficult in the script, such as changing the username and password, so such requirements were not included in the grading rubric for the script. The most substantial additional requirement for the script was that it uses “standard and stable external resources whenever possible (e.g., [www.vagrantbox.es](http://www.vagrantbox.es), rather than a hand-built box)”.

#### 4.2.6 *Redistributing the Tools*

Teams were required to establish GitHub repositories for each tool that their team was assigned – working or not. The main goal was to share the work the teams had done with others. We created an organization to house each repository, which can be found online.<sup>7</sup>

If the original tool contained a license to redistribute the source code, or the paper authors gave us explicit permission to do so, we redistributed that code in our tool’s repository. The same applied to the tools’ binary. If the tool was working, vagrant scripts were also included in the repository, regardless of whether we had permission to redistribute the tool; if we did not, users would have to contact the original authors and drop in the tool (a binary, for instance) for the script to work. Even if the tool was not available to us, we created a mostly empty repository, for consistency.

For tools for which we had source code, whenever possible teams included the history of that source code, for completeness. Ideally, tools that were already hosted on GitHub could be forked by the team, maintaining not just history but also an explicit connection to the original repository. When original tool were hosted elsewhere, such as in subversion on Google Projects, teams migrated source code history to GitHub.

Teams added a readme file to each tool’s repository that conveyed some basic information about the tool, including

links to the original paper and original project webpage, as well as acknowledgements. The rubric outlined a number of small, specific requirements for the readme to ensure consistency.<sup>8</sup>

When teams had permission to redistribute the virtual machine image containing the tool, readmes also contained links to those images. The images were hosted on the instructor’s Google Drive account, which provides unlimited cloud storage.<sup>9</sup> Such large capacity storage is necessary because each image occupies several gigabytes of space.

#### 4.2.7 *Presenting Tools*

For the communication part of the class, teams were to present each working tool in front of the class for a 5 minute demonstration. The grading rubric specified, among other requirements, that the presentations explained what problem the tool was built to solve, to give a simple enough example that the tool could be understood, and to be realistic enough to be compelling.<sup>10</sup> Teams also prepared a video demo of each working tool, posted on YouTube and linked to from the virtual machine image and the GitHub readme.

#### 4.2.8 *Evaluations*

Teams’ deliverables were evaluated by their peers and by two teachers’ assistants (TAs). While both peers and TAs evaluated the deliverables for quality and consistency, only the TA evaluations counted towards teams’ grades.

Teams completed peer evaluations of other teams’ hand-built virtual machines and GitHub repositories. This entailed reading the original paper to understand how the tool was supposed to work, using the tool in the virtual machine, and comparing deliverables to the grading rubrics. When students found defects, they filed issues in the repositories issue tracker and provided fixes for simple defects using pull requests.

Hand-built virtual machine images were evaluated in two rounds by the TAs, where students enhanced their images between rounds and the rubrics were updated for the second round based on what the TAs observed in the first. Likewise, vagrant scripts were evaluated in two rounds by the TAs. Finally, TAs evaluated all artifacts (images, vagrant scripts, and repositories) together in a final project evaluation.

#### 4.2.9 *Data Collection and Contributions*

As part of the final project submission, the instructor collected data about each tool in the form of a survey. This tool survey asked, for instance, how long the team estimated they spent on getting the tool working. Each student also submitted a second survey that asked students what they thought about the project, and whether they wanted to participate as an author of this paper. Data was also collected from standardized and anonymous course evaluation forms, forms that are used for all courses across the university.

This paper was written primarily by the first author, the instructor of the course, after the course had ended. The

<sup>5</sup>[https://docs.google.com/spreadsheets/d/1zGwF\\_TbGCMrwH1vqY2\\_OHAQg7umd4zyAepHiirdb-yM/edit?usp=sharing](https://docs.google.com/spreadsheets/d/1zGwF_TbGCMrwH1vqY2_OHAQg7umd4zyAepHiirdb-yM/edit?usp=sharing)

<sup>6</sup>[https://docs.google.com/spreadsheets/d/1qSt63AjwyiCRHCU-MMoTSDx\\_qPuWn7U4\\_QmjwC\\_Ghds/edit?usp=sharing](https://docs.google.com/spreadsheets/d/1qSt63AjwyiCRHCU-MMoTSDx_qPuWn7U4_QmjwC_Ghds/edit?usp=sharing)

<sup>7</sup><https://github.com/SoftwareEngineeringToolDemos/>

<sup>8</sup>[https://docs.google.com/spreadsheets/d/1ufh-XiNFY1jRvehDN\\_JuLYQTRVHzuEDru6stPAyxdc/edit?usp=sharing](https://docs.google.com/spreadsheets/d/1ufh-XiNFY1jRvehDN_JuLYQTRVHzuEDru6stPAyxdc/edit?usp=sharing)

<sup>9</sup><https://support.google.com/a/answer/2856827>

<sup>10</sup>[https://docs.google.com/spreadsheets/d/1Mb8yacxIP233kmjn3J\\_JtJfOfuJlJmSHnMKtr1CVxCM/edit?usp=sharing](https://docs.google.com/spreadsheets/d/1Mb8yacxIP233kmjn3J_JtJfOfuJlJmSHnMKtr1CVxCM/edit?usp=sharing)

paper’s source and history can also be found on GitHub.<sup>11</sup>  
The technical work was con

## 5. EXAMPLE TOOL

Sketch example tool here, with all the pieces. It would be nice if it exemplified some issues brought up later, such as technology stack licensing.

## 6. PROCEDURE

No cost, but demos ok. (How many did this happen with?)  
Tools where tool could conceivably be put into a virtual machine. How many didn’t fit?

## 7. RESULTS

### 7.1 RQ1 and RQ2

#### 7.1.1 Effort

Average person-hours per tool, distribution (min, max, box plot)  
Time invested in evaluating non-working tools.

#### 7.1.2 Challenge: Tool Cannot Be Obtained

What percent of tool links were dead? What percent of tools said they were available in the paper, but the tool could not be obtained? What percent of tools were being planned to be commercialized? What were actually commercial? What percent of tools could we use if we had paid for them?

#### 7.1.3 Challenge: Disappearing Tools

One author said tool just doesn’t exist anymore.  
Another author had to dig tool out of long term archive.  
Several authors (what percent?) has tools hosted on Google code, even though it was dying. Did we save ’em?

#### 7.1.4 Challenge: Author non-responsive

#### 7.1.5 Challenge: Technical Difficulties

#### 7.1.6 Challenge: Inconsistencies

One tool required to different VMs because two features needed different prereqs.  
Some tools versioned differently.

### 7.2 RQ3 and RQ4

#### 7.2.1 Challenge: Tool Licensing

University grey area.

#### 7.2.2 Challenge: Technology Stack Licensing

#### 7.2.3 Challenge: Author Doesn’t Want Redistribution

Even when tool is available

## 8. CONCLUSIONS

Some other things.

## 9. ADDITIONAL AUTHORS

Additional authors: Shabbir Abdul (sabdul@ncsu.edu), Varun Aettapu (vaettap@ncsu.edu), Sumeet Agarwal (sagarwa6@ncsu.edu), Sindhu Anangur Vairavel (sanangu@ncsu.edu), Rishi Avinash Anne (raanne@ncsu.edu), Haris Mahmood Ansari (hmansari@ncsu.edu), Ankit Bhandari (abhanda3@ncsu.edu), Anand Bhanu (bhanua@ncsu.edu), Aditya Vinayak Bhise (avbhise@ncsu.edu), Saikrishna Teja Bobba (sbobba3@ncsu.edu), Vineela Boddula (vboddul@ncsu.edu), Venkata Krishna Sailesh Bommiseti (vbommis@ncsu.edu), Dwayne Christian (Chris) Brown (dcbrow10@ncsu.edu), Peter Morgan Chen (pmchen@ncsu.edu), Yi Chun (Yi-Chun) Chen (ychen74@ncsu.edu), Nikhil Chinthapalle (nchinth@ncsu.edu), Karan Singh Dagar (kdagar@ncsu.edu), Joseph Decker (jdecke@ncsu.edu), Pankti Rakeshkumar Desai (prdesai2@ncsu.edu), Jayant Dhawan (jdhawan2@ncsu.edu), Yihuan Dong (ydong2@ncsu.edu), Sarah Elizabeth Elder (seelder@ncsu.edu), Shrenuj Gunvant Gandhi (sgandhi4@ncsu.edu), Jennifer Michelle Green (jmgree17@ncsu.edu), Mohammed Hasibul Hassan (mhassan@ncsu.edu), Satish Inampudi (sinampu@ncsu.edu), Pragyan Paramita Jena (ppjena@ncsu.edu), Bhargav Rahul (Bhargav) Jhaveri (bjhaver@ncsu.edu), Apoorv Vijay Joshi (avjoshi@ncsu.edu), Nikhil Josyabhatla (njosyab@ncsu.edu), Sujith Katakam (skatakam@ncsu.edu), Juzer Husainy Khambaty (jkhkamba@ncsu.edu), Aneesh Arvind Kher (aakher@ncsu.edu), Craig Kimpel (ckimpal@ncsu.edu), Siddhartha Kollipara (skollip@ncsu.edu), Asish Prabhakar Kot-tala (akottal@ncsu.edu), Abishek Kumar (akumar21@ncsu.edu), Harini Reddy Kumbum (hkumbum@ncsu.edu), Nitish Pradeep Limaye (nplimaye@ncsu.edu), Apoorv Mahajan (amahaja3@ncsu.edu), Sai Sindhur Malleni (smallen3@ncsu.edu), Sudha Manchukonda (smanchu@ncsu.edu), Kavita Maral Mehta (kmmeha@ncsu.edu), Justin Alan Middleton (jamidl2@ncsu.edu), Ramakant Moka (rmoka@ncsu.edu), Eesha Gopalakrishna Mulky (egmulky@ncsu.edu), Gauri Naik (gnaik2@ncsu.edu), Shraddha Anil Naik (sanaik2@ncsu.edu), Yashwanth Nallabothu (ynallab@ncsu.edu), Yogesh Nandakumar (ynandak@ncsu.edu), Kairav Sai Padarth (kspadart@ncsu.edu), Pulkesh Kumar Yadav Pannalal (ppannal@ncsu.edu), Sattwik Pati (spati2@ncsu.edu), Kahan Prabhu (kprabhu@ncsu.edu), Shashank Goud Pulimamidi (spulima@ncsu.edu), Gargi Sandeep Rajadhyaksha (gsrajadh@ncsu.edu), Priyadarshini Rajagopal (prajago4@ncsu.edu), Venkatesh Sambandamoorthy (vsamban@ncsu.edu), Mohan Sammeta (msammet@ncsu.edu), Shaown Sarker (ssarker@ncsu.edu), Anshita Sayal (asayal@ncsu.edu), Vrushti Kamleshkumar Shah (vkshah@ncsu.edu), Esha Sharma (esharma2@ncsu.edu), Saurav Shekhar (sshekha3@ncsu.edu), Sarthak Prabhakar Shetty (spshetty@ncsu.edu), Manish Ramashankar Singh (mrsingh@ncsu.edu), Ankush Kumar Singh (asingh21@ncsu.edu), Vinay Kumar Suryadevara (vksuryad@ncsu.edu), Sumit Kumar Tomer (sktomer@ncsu.edu), Akriti Tripathi (atripat4@ncsu.edu), Jennifer Tsan (jtsan@ncsu.edu), Vivekananda Vakkalanka (vvakkal@ncsu.edu), Alexander Valkovsky (avalkov@ncsu.edu), Rishi Kumar Vardhini (rkvardhi@ncsu.edu) and Manav Verma (mverma4@ncsu.edu).

## 10. REFERENCES

- [1] C. Collberg, T. Proebsting, and A. M. Warren. Repeatability and benefaction in computer systems research. Technical report, University of Arizona, 2015. TR 14-04.

<sup>11</sup><https://github.com/SoftwareEngineeringToolDemos/paper>

- [2] C. Collberg and T. A. Proebsting. Repeatability in computer systems research. *Communications of the ACM*, 59(3):62–69, 2016.
- [3] A. Desai, S. Gulwani, V. Hingorani, N. Jain, A. Karkare, M. Marron, S. R., and S. Roy. Program synthesis using natural language. In L. K. Dillon, W. Visser, and L. Williams, editors, *Proceedings of the 38th International Conference on Software Engineering, ICSE 2016, Austin, TX, USA, May 14–22, 2016*, pages 345–356. ACM, 2016.
- [4] M. D. Ernst, A. Lovato, D. Macedonio, F. Spoto, and J. Thaine. Locking discipline inference and checking. In *ICSE’16, Proceedings of the 38th International Conference on Software Engineering*, Austin, TX, USA, May 18–20, 2016.
- [5] J. M. González-Barahona and G. Robles. On the reproducibility of empirical software engineering studies based on data retrieved from development repositories. *Empirical Software Engineering*, 17(1-2):75–89, 2012.
- [6] D. C. Ince, L. Hatton, and J. Graham-Cumming. The case for open computer programs. *Nature*, 482(7386):485–488, 2012.
- [7] C. Klein, J. Clements, C. Dimoulas, C. Eastlund, M. Felleisen, M. Flatt, J. A. McCarthy, J. Raskind, S. Tobin-Hochstadt, and R. B. Findler. Run your research: on the effectiveness of lightweight mechanization. In *ACM SIGPLAN Notices*, volume 47, pages 285–296. ACM, 2012.
- [8] J. Kovacevic. How to encourage and publish reproducible research. In *Acoustics, Speech and Signal Processing, 2007. ICASSP 2007. IEEE International Conference on*, volume 4, pages IV–1273. IEEE, 2007.
- [9] X. Lu, X. Liu, H. Li, T. Xie, Q. Mei, D. Hao, G. Huang, and F. Feng. Prada: prioritizing android devices for apps by mining large-scale usage data. In *Proceedings of the 38th International Conference on Software Engineering*, pages 3–13. ACM, 2016.
- [10] A. Meneely, L. Williams, and E. F. Gehring. Rose: a repository of education-friendly open-source projects. In *ACM SIGCSE Bulletin*, volume 40, pages 7–11. ACM, 2008.
- [11] M. Pedroni, T. Bay, M. Oriol, and A. Pedroni. Open source projects in programming courses. *ACM SIGCSE Bulletin*, 39(1):454–458, 2007.
- [12] V. Stodden. Enabling reproducible research: Licensing for scientific innovation. *Int’l J. Comm. L. & Pol’y*, 13:1, 2009.
- [13] P. Vandewalle, J. Kovačević, and M. Vetterli. Reproducible research in signal processing. *Signal Processing Magazine, IEEE*, 26(3):37–47, 2009.